

Conflict-Aware Cross-View Disaster Triage Across Spatial Observation Regimes [experiments]

Yifan Yang

Department of Geography
Texas A&M University
College Station, Texas, USA

Lei Zou

Department of Geography
Texas A&M University
College Station, Texas, USA

ABSTRACT

Cross-view disaster assessment combines ground-level and overhead imagery to infer building damage, but aggregate F1 hides examples where the two views disagree. This short paper introduces Conflict-Aware Evaluation (CAE), a reporting protocol that measures overall F1, conflict-subset F1, cross-view gain over the stronger single-view baseline, and the direction in which one view corrects the other. We apply CAE to an Eaton/Altadena wildfire dataset with property-centered ground imagery and to the Hurricane Ian CVDIsaster benchmark with panoramic ground imagery. The main result is deliberately mixed. On the wildfire sensitivity split, cross-view fusion improves conflict F1 by 0.0388. On the leakage-controlled hurricane endpoint split, the gain becomes negative ($\Delta_{\text{view}} = -0.0153$), although aggregate F1 makes cross-view and street-only models look nearly tied. Directional and alignment diagnostics show that fusion helps when the ground view observes the target building, but can lose value when the ground view becomes panoramic context. CAE therefore turns paired-view triage into a stress test for when fusion helps, fails, or should defer.

KEYWORDS

disaster assessment, cross-view learning, multimodal fusion, GeoAI, conflict-aware evaluation

1 INTRODUCTION

Post-disaster building assessment increasingly uses several visual viewpoints. Overhead imagery gives broad coverage, while ground imagery can expose facade condition, debris, and local structural cues. Cross-view models seek to combine these signals, but average test scores hide the cases that matter most for deployment: buildings where the views disagree.

This paper treats paired-view triage as conflict resolution. Each example has a ground image, an overhead patch, and a binary damage label. We ask whether fusion improves the subset where the single-view predictors disagree, rather than whether it only improves average F1. This framing matters because a property inspection photograph and a panoramic street scene are both ground views, yet they observe the labeled building in different ways.

The work builds on overhead damage benchmarks [2], fusion models for disaster assessment [5], ground-view damage studies [6–8], and the paired Hurricane Ian CVDIsaster benchmark [4]. CAE is related to selective and disagreement-aware evaluation [1], but its reporting unit is spatial: the disagreement subset, the single-view correction direction, and the gain or loss from fusion under a given observation regime. The paper makes three claims: CAE exposes conflict behavior, leakage-controlled hurricane results set a

Table 1: Evaluation settings, sample counts, and leakage controls.

Setting	Positive class	Ground regime	N	Role
WF sens.	affected+	property	19,720	main CAE
H end.	severe	panoramic	2,583	diagnostic
H mod.+sev.	moderate+	panoramic	4,121	sensitivity
H clean	severe	panoramic	2,583	leakage control

negative boundary for fusion, and spatial observation regime helps explain the dominant-view switch.

2 DATA AND PROTOCOL

2.1 Paired-View Triage Settings

Let $\mathcal{D} = \{(x_i^{(s)}, x_i^{(r)}, y_i)\}_{i=1}^N$ be a paired disaster triage dataset with a ground image, an overhead patch, and a binary label. We compare street-only (S), remote-only (R), and cross-view (CV) predictors. CV uses late fusion of ResNet18 features [3]: both embeddings, their absolute difference, and their elementwise product.

Table 1 summarizes the four evaluation settings. The wildfire setting uses 19,720 usable Eaton/Altadena paired records derived from property-level inspection photographs and matched overhead crops. The source index retains object identifiers, coordinates, crop boxes, overhead tile IDs, and coverage metadata, enabling object-level grouping and tile-level spatial aggregation. The positive class is Affected + Minor + Major + Destroyed. The hurricane setting uses the Hurricane Ian CVDIsaster benchmark [4], with 4,121 paired samples across minor, moderate, and severe damage. We evaluate both the endpoint task (minor versus severe) and a broader moderate-plus-severe task. The comparison is not a causal test of wildfire versus hurricane; it contrasts property-centric and panoramic observation regimes.

2.2 Conflict-Aware Evaluation

For each example, let $p_i^{(s)}$ and $p_i^{(r)}$ denote the positive-class probabilities predicted by the street-only and remote-only models. CAE defines the conflict subset

$$C_\tau = \{i : |p_i^{(s)} - p_i^{(r)}| > \tau\}, \quad (1)$$

with $\tau = 0.1$ in the main analysis. For each model m , CAE reports standard overall F1 and conflict F1 on C_τ . It then computes the cross-view conflict gain

$$\Delta_{\text{view}} = \text{F1}(\text{crossview}; C_\tau) - \max\{\text{F1}(\text{street}; C_\tau), \text{F1}(\text{remote}; C_\tau)\}. \quad (2)$$

The compact CAE report for each model is therefore

$$\mathbf{r}_\tau(m) = (\text{F1}(m; \mathcal{D}), \text{F1}(m; C_\tau), n_c/N), \quad n_c = |C_\tau|. \quad (3)$$




Figure 1: Study design and spatial observation regimes. The wildfire setting uses property-centered inspection imagery paired with overhead patches, whereas the hurricane setting uses panoramic ground imagery paired with overhead patches. This visual difference motivates the regime-aware interpretation tested by CAE.




Figure 2: Conflict-Aware Evaluation (CAE) protocol. CAE isolates examples where street-only and remote-only predictors disagree, then reports whether cross-view fusion improves conflict-subset F1 over the stronger single-view baseline. Directional arbitration links model behavior to spatial observation regime.

We also decompose conflicts by correction direction. Let $\hat{y}_i^{(s)} = 1\{p_i^{(s)} \geq 0.5\}$ and define $\hat{y}_i^{(r)}$ analogously. With $c_i^{(v)} = 1\{\hat{y}_i^{(v)} = y_i\}$, the two arbitration rates are

$$\rho_{S \rightarrow O} = \frac{1}{n_c} \sum_{i \in C_\tau} c_i^{(s)} (1 - c_i^{(r)}), \quad (4)$$

$$\rho_{O \rightarrow S} = \frac{1}{n_c} \sum_{i \in C_\tau} c_i^{(r)} (1 - c_i^{(s)}). \quad (5)$$

$S \rightarrow O$ denotes street-correct and remote-wrong cases, and $O \rightarrow S$ denotes the reverse. Because a prior hurricane endpoint split had object-level leakage, the grouped-clean endpoint split is the conservative control.

2.3 Implementation Details

Images are resized to 224 by 224 pixels and normalized for ImageNet-pretrained ResNet18 encoders. Each view embedding is projected to 256 dimensions. Models use AdamW with learning rate 10^{-4} for 10 epochs, and the best-validation-F1 checkpoint is evaluated on the held-out test split.

For reproducibility, the wildfire release is organized around split manifests and a data-card summary rather than raw-image redistribution. The manifests keep object IDs, image-source fields, coordinates, matched overhead tile names, binary labels, and filtering

flags. This makes the grouped split auditable and lets future work reconstruct the paired-view benchmark from the original imagery sources when access permits.

3 RESULTS

Table 2 is the CAE scorecard. Average F1 alone would make CV look mostly helpful or tied. Conflict F1 changes the reading. Wildfire has a clear positive conflict gain (+0.0388), while the leakage-controlled hurricane split has a negative gain (-0.0153) even though S and CV are nearly tied overall.

Figure 3 explains why the mixed gains are useful rather than inconvenient. In wildfire, $S \rightarrow O$ conflicts are more common than $O \rightarrow S$ conflicts, so property-centered ground imagery is the stronger arbitrator. In the broader hurricane split, the ordering reverses, which is consistent with panoramic views that contain more scene context and weaker target alignment. The threshold sweep and alignment diagnostics show that the CAE pattern is not only a single-threshold artifact.

Additional checks support the same ordering without changing the claim. At $\tau = \{0.1, 0.3, 0.5\}$, CV remains strongest on wildfire conflict accuracy (0.8349, 0.7821, 0.7538). At $\tau = 0.1$, wildfire conflict accuracy is 0.7885 for CV, 0.7356 for S, and 0.6529 for R; the hurricane moderate-plus-severe split gives 0.7500, 0.7031, and




Figure 3: Result diagnostics. (a) Directional conflicts show the dominant-view switch. (b) Wildfire CV remains strongest across thresholds. (c–d) Alignment diagnostics support the regime interpretation. (e) Tile-level maps localize where fusion helps or fails.

Table 2: Conflict-Aware Evaluation summary. Conflict F1 uses C_τ with $\tau = 0.1$. WF is wildfire; H is hurricane. Gain is reported only for cross-view rows.

Setting	M	All	C-F1	Gain	n_c
WF sens.	S	0.9392	0.6230	–	435
WF sens.	R	0.9238	0.5237	–	435
WF sens.	CV	0.9473	0.6618	+0.0388	435
H end.	S	0.9055	0.6842	–	32
H end.	R	0.8969	0.5806	–	32
H end.	CV	0.9192	0.7647	+0.0805	32
H mod.+sev.	S	0.8738	0.7324	–	128
H mod.+sev.	R	0.8873	0.7742	–	128
H mod.+sev.	CV	0.8915	0.7867	+0.0125	128
H clean	S	0.9016	0.8101	–	80
H clean	R	0.8385	0.6526	–	80
H clean	CV	0.9008	0.7949	-0.0153	80

0.7266. Multi-seed validation also keeps CV slightly ahead in base F1, while generic CLIP and DINOv2 features underperform the supervised CNN baselines in this disaster-specific setting.

The tile-level panel adds a geographic check on this interpretation. Wildfire damage, disagreement density, and cross-view gain are spatially uneven across overhead tiles. High damage rate does not automatically imply high conflict, and high conflict does not automatically imply positive CV gain. This matters for disaster response because deployment errors are rarely spread uniformly over a region. A model can look strong on a pooled metric while failing in a cluster of parcels with unusual viewing angle, smoke

shadow, vegetation cover, or partial structure visibility. CAE therefore works as both a model-comparison metric and a spatial audit layer: it identifies where the disagreement cases are, which view wins there, and whether the fused model resolves or amplifies the local ambiguity.

4 DISCUSSION

The grouped-clean hurricane result is the key stress test: more visual evidence does not always improve decisions once object-level leakage is controlled. CAE therefore exposes a deployment boundary, not only a performance gain.

The conflict subset is not meant to replace the full test set. It is a second readout that asks a narrower deployment question: when the two available views would support different decisions, does fusion add useful evidence or blur the better view? This differs from selecting examples by loss or label difficulty after evaluating CV. The gate is defined only by the two single-view probability outputs, so it can be computed before choosing whether to trust fusion. The directional counts then make the error source visible. If most conflicts are $S \rightarrow O$, the street view supplies the missing local cue; if most are $O \rightarrow S$, the overhead view is the stronger arbitrator. A useful paired-view benchmark should report both cases because a positive average gain can hide a failure on the examples where fusion is expected to help.

Spatial observation regime makes this boundary plausible. In the wildfire records, the photographer usually stands near the parcel and aims at the structure that receives the label. The ground image can therefore carry direct evidence about charred walls, damaged openings, debris against the facade, or an intact exterior. In the hurricane records, the panoramic image often captures the street segment rather than a single labeled footprint. It may show flooded roads, vegetation loss, roof debris, and neighboring buildings, but the target building can be small, peripheral, or occluded. The mechanism is therefore how each view observes the target, not simply “street versus satellite.” The same architecture sees two views in both settings, but the semantic access of each view differs.

Operationally, cases where S and R agree may use the agreement decision; cases where they disagree should route to CV or human review. Future datasets should collect property-centered and panoramic views for the same target buildings to test target alignment directly.

The threshold τ should be interpreted as an operating knob rather than a fixed scientific constant. A small threshold sends more examples into the conflict set and measures broad disagreement. A larger threshold keeps only the strongest single-view contradictions, which may be more suitable when human review capacity is limited. The main analysis uses $\tau = 0.1$ because it leaves enough examples for stable reporting while still isolating visible disagreement. The threshold sweep in Figure 3 is therefore part of the claim. It shows whether the conclusion survives when the triage system keeps fewer but harder cases.

The study still has limits. The wildfire and hurricane settings differ in hazard type, geography, collection protocol, and label construction, so the paper treats observation regime as an explanatory lens rather than a causal proof. The compact image-alignment proxies in Figure 3 are diagnostic, not a replacement for field-verified viewpoint geometry. A stronger benchmark would collect several ground-view regimes for the same buildings and report CAE under identical spatial splits.

Two implementation details also deserve care in future work. First, the fusion model is deliberately simple, so the negative grouped-clean hurricane result should not be read as a universal limit of cross-view learning. It shows that late fusion can fail to improve the hardest cases after leakage control. Second, object grouping matters as much as architecture. If visually similar views of the same object appear across splits, a model can appear to learn cross-view reasoning while partly memorizing object-specific cues. CAE is most useful when paired with grouped splits, transparent manifests, and geographic diagnostics. Those controls turn fusion from a single leaderboard number into an auditable decision rule.

5 CONCLUSION

CAE reframes paired-view disaster triage as a conflict-resolution problem. It shows a positive wildfire conflict gain, a negative grouped-clean hurricane gain, and a dominant-view switch tied to spatial observation regime. Disaster benchmarks should therefore report conflict-subset behavior, directional arbitration, and threshold sensitivity alongside aggregate scores. This added reporting does not require a new model. It changes what the evaluation asks: whether fusion helps exactly where single-view evidence disagrees.

The recommended reporting template is compact. For each split, authors should publish full-test F1, conflict F1 for each view and the fused model, the number of conflict cases, the cross-view gain over the best single view, and the $S \rightarrow O$ versus $O \rightarrow S$ correction balance. When possible, these numbers should be paired with grouped splits and a small spatial summary of where conflicts occur. This template would make negative results useful rather than embarrassing. If fusion loses on the conflict subset, the benchmark has revealed a boundary condition. If fusion wins, the directional counts show which view supplied the missing evidence. Either outcome gives disaster analysts a more actionable evaluation than a single average score.

For this reason, the contribution is evaluative rather than architectural. The paper does not argue that one fusion block is optimal, or that one disaster type always favors one view. It argues that paired-view benchmarks need a sharper unit of analysis. Conflict cases reveal the operational value of fusion because they are the cases where two plausible sensors disagree. Reporting those cases also makes future architectural work easier to judge: a stronger model should not only improve pooled F1, but also reduce the unresolved conflict left by its single-view components.

ACKNOWLEDGMENTS

The authors used AI-based writing assistance for grammar checking and editorial compression, reviewed all content, and take full responsibility for the work. Supported by the Texas A&M University Environment and Sustainability Initiative (ESI) through the Environment and Sustainability Graduate Fellow Award.

REFERENCES

- [1] Yonatan Geifman and Ran El-Yaniv. 2017. Selective Classification for Deep Neural Networks. In *Advances in Neural Information Processing Systems*, Vol. 30. 4878–4887.
- [2] Ritwik Gupta, Bryce Goodman, Nirav Patel, Ricky Hosfelt, Sandra Sajeev, Eric Heim, Jigar Doshi, Keane Lucas, Howie Choset, and Matthew Gaston. 2019. Creating xBD: A Dataset for Assessing Building Damage from Satellite Imagery. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops*. 10–17.
- [3] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. Deep Residual Learning for Image Recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 770–778.
- [4] Hao Li, Fabian Deuser, Wenping Yin, Xuanshu Luo, Paul Walther, Gengchen Mai, Wei Huang, and Martin Werner. 2025. Cross-View Geolocalization and Disaster Mapping with Street-View and VHR Satellite Imagery: A Case Study of Hurricane IAN. *ISPRS Journal of Photogrammetry and Remote Sensing* 220 (2025), 841–854.
- [5] Weiwei Xiao, Jingyong Su, Yongyong Chen, and Guofeng Cao. 2023. Cross-Scale-Guided Fusion Transformer for Disaster Assessment Using Satellite Imagery. *IEEE Transactions on Geoscience and Remote Sensing* 61 (2023), 1–12.
- [6] Zhuoqun Xue, Xiaojian Zhang, David O Prevatt, Jennifer Bridge, Susu Xu, and Xilei Zhao. 2024. Post-Hurricane Building Damage Assessment Using Street-View Imagery and Structured Data: A Multi-Modal Deep Learning Approach. *arXiv preprint arXiv:2404.07399* (2024).
- [7] Yifan Yang, Lei Zou, Bing Zhou, Daoyang Li, Binbin Lin, Joynal Abedin, and Mingzheng Yang. 2025. Hyperlocal Disaster Damage Assessment Using Bi-Temporal Street-View Imagery and Pre-Trained Vision Models. *Computers, Environment and Urban Systems* 121 (2025), 102335.
- [8] Wei Zhai and Zhong-Ren Peng. 2020. Damage Assessment Using Google Street View: Evidence from Hurricane Michael in Mexico Beach, Florida. *Applied Geography* 123 (2020), 102252.